

Past Exams

- $w^* = \operatorname{argmin}_{w \in \mathbb{R}^D} \|y - Xw\|_2^2$ is unique only if X is full rank
- Find minimum of function wrt a parameter θ by taking its gradient wrt θ , set it to 0
- Probability likelihood $\prod$; Log-likelihood $\sum$; $w^T w = \|w\|_2^2$
- All linear/affine functions, sum of convex functions and composition of linear/affine functions and convex functions are convex.

Optimization

- Convex: $\mathcal{L}(\lambda w + (1 - \lambda)w') \leq \lambda \mathcal{L}(w) + (1 - \lambda)\mathcal{L}(w') \forall \lambda \in [0, 1]$ or Hessian is PSD.
- Loss function should be symmetrical and penalize large and very large mistakes similarly (so outliers don't mess up the model, as we cannot just remove them)
- GD: $w^{(t+1)} = w^{(t)} - \gamma \nabla \mathcal{L}(w^{(t)})$ (γ step-size; for L -smooth convex $\mathcal{L}$, convergence for $0 < \gamma < 2/L$; linear models: $O(ND)/\text{iter}$).
- SGD: pick random n_t , $g_t = \nabla \mathcal{L}_{n_t}(w^{(t)})$, unbiased $\mathbb{E}[g_t] = \nabla \mathcal{L}(w^{(t)})$ ($O(D)/\text{step}$)
- Mini-batch: lower variance, parallelizable; $g_t = \frac{1}{|B|} \sum_{n \in B} \nabla \mathcal{L}_n(w^{(t)})$
- Momentum: $m_{t+1} = \beta m_t + (1 - \beta)g_t$, $w_{t+1} = w_t - \gamma m_{t+1}$ (smooths noisy SGD, reduces oscillations, momentum = moving average of the gradient).
- Adam: $m_{t+1} = \beta_1 m_t + (1 - \beta_1)g_t$, $v_{t+1} = \beta_2 v_t + (1 - \beta_2)g_t \odot g_t$ (momentum + adaptive per-coordinate scaling; more stable steps; forgets older weights faster; scale-invariant to gradient). $\hat{m} = \frac{m}{1 - \beta_1^{t+1}}$, $\hat{v} = \frac{v}{1 - \beta_2^{t+1}}$; $w_{t+1} = w_t - \gamma \hat{m}/(\sqrt{\hat{v}} + \epsilon)$.
- SignSGD: $w_{t+1} = w_t - \gamma \operatorname{sign}(g_t)$ (1 bit/coord; distributed sys; convergence issues)
- Subgradient g : $f(y) \geq f(x) + g^T(y - x)$, $\forall y$, any slope of a line that touches f at a point but stays entirely below f 's curve. Used in nonsmooth convex GD as gradient
- Projected GD (constraints $w \in C$ convex): $w^{(t+1)} = P_C(w^{(t)} - \gamma \nabla \mathcal{L}(w^{(t)}))$ (brick wall $I_C(w) = 0$ if $w \in C$, 0 otherwise; penalty method $\operatorname{argmin} \mathcal{L}(w) + \gamma \|Aw - b\|_2^2$).
- Feature stdization (pre-conditioning) helps ill-conditioning (features w/ diff. scales)
- GD iterations are cheaper than Newton's method, but require more to converge

Linear Regression

- Linear Models are scale & shift invariant when there is no regularization
- MSE (Quadratic Loss): $L(w) = \frac{1}{2N} \|y - Xw\|_2^2 = \frac{1}{2N} \sum_{n=1}^N (y_n - x_n^T w)^2$.
- $\nabla L(w) = -\frac{1}{N} X^T (y - Xw)$. Hessian: $\nabla^2 L(w) = \frac{1}{N} X^T X$.
- Closed-form Least Squares (full col rank): $w^* = (X^T X)^{-1} X^T y$; if rank-deficient use regularization/pseudo-inverse $\text{np.linalg.solve}(\text{Gram } X^T X: O(ND^2), \text{solve}: O(D^3))$.

Overfitting & Underfitting

- Underfit (high bias): train err high & test err high (model too simple/bad features).
- Overfit (high variance): train err low but test err high (model too complex).
- Polynomial features: $\phi(x) = [1, x, x^2, \dots, x^M]$, linear in ϕ but nonlinear in x .
- Regularization: $\min_w \mathcal{L}(w) + \Omega(w)$ convex; promotes sparsity; optimum at corners, equivalent to imposing a Laplace prior on the weights, diamond-shaped
- Log-likelihood (Gaussian): $y_n = x_n^T w + \epsilon_n$, $\epsilon_n \sim \mathcal{N}(0, \sigma^2)$, $\log p(y | X, w) = \log \prod_{n=1}^N p(y_n | x_n, w) = \log \prod_{n=1}^N \mathcal{N}(y_n | x_n^T w, \sigma^2) = -\frac{1}{2\sigma^2} \sum_{n=1}^N (y_n - x_n^T w)^2 + c$
- Log-likelihood (Laplace): uses MAE instead, $p(y_n | x_n, w) = \frac{1}{2b} e^{-\frac{1}{b} |y_n - x_n^T w|}$.

Generalization, Model Selection, and Validation

- Generalization/True Error/Risk: $\mathcal{L}_{\mathcal{D}}(f) = \mathbb{E}_{(x,y) \sim \mathcal{D}} [\ell(y, f(x))]$, the expected error over $\mathcal{D}$ true real-world distribution. Generalization gap: $|\mathcal{L}_{\mathcal{D}}(f) - \mathcal{L}_S(f)|$
- Empirical Error: $\mathcal{L}_S(f) = \frac{1}{|\mathcal{S}|} \sum_{(x_n, y_n) \in \mathcal{S}} \ell(y_n, f(x_n))$, unbiased estimator of the true error $\mathcal{L}_{\mathcal{D}}(f)$ when having enough data (Law of Large Numbers) Generalization gap
- Generalization Gap with test set: δ confidence threshold, $\mathbb{P} \left[|\mathcal{L}_{\mathcal{D}}(f) - \mathcal{L}_{S_{\text{test}}}(f)| \geq \sqrt{\frac{(b-a)^2 \ln(2/\delta)}{2|S_{\text{test}}|}} \right] \leq \delta$; gap = $O(1/\sqrt{|S_{\text{test}}|})$.
- Probabilistic up. bound of test err $\mathbb{P}[\mathcal{L}_{\mathcal{D}}(f) \geq \mathcal{L}_{S_{\text{test}}}(f) + \sqrt{\frac{(b-a)^2 \ln(2/\delta)}{2|S_{\text{test}}|}}] \leq \delta$
- Hoeffding Inequality: $\mathbb{P} \left[\left| \frac{1}{N} \sum_{n=1}^N \Theta_n - \mathbb{E}[\Theta] \right| \geq \epsilon \right] \leq 2e^{-2N\epsilon^2/(b-a)^2}$, for i.i.d. $\Theta_n \in [a, b]$, $\Theta_i = \ell(y_i, f(x_i))$, empirical mean concentrated around the true one.
- K models (hyperparameter search): $\mathbb{P}[\max_{k \leq K} |\mathcal{L}_{\mathcal{D}}(f_k) - \mathcal{L}_{S_{\text{test}}}(f_k)| \geq \sqrt{\frac{(b-a)^2 \ln(2K/\delta)}{2|S_{\text{test}}|}}] \leq \delta$, can test many models without diverging a lot.
- Classification
- 0-1 loss (binary): $\ell(y, f(x)) = \mathbf{1}_{y \neq f(x)}$; nonconvex/discontinuous.
- Bayes classifier $f^*(x) \in \operatorname{argmax}_{y \in \{-1, 1\}} \mathbb{P}(Y = y | X = x)$, minimize 0-1 loss in $\mathcal{D}$
- ERM (0-1): $\min_f \frac{1}{N} \sum_{n=1}^N \mathbf{1}_{f(x_n) \neq y_n} = \min_g \frac{1}{N} \sum_{n=1}^N \mathbf{1}_{y_n g(x_n) \leq 0}$ with $f(x) = \operatorname{sign}(g(x))$, g model output.
- Convex surrogate: replace $\mathbf{1}_{y_n g(x_n) \leq 0}$ by convex upper bound $\phi(y_n g(x_n))$ (e.g. hinge or logistic loss) $\Rightarrow$ convex ERM.
- Over-parametrization ($N < D$): if separable data, GD classification $\approx$ GD regression.

Logistic Regression

- Sigmoid $\sigma(\eta) = \frac{1}{1+e^{-\eta}} = \frac{e^{\eta}}{1+e^{\eta}}$, $\sigma'(\eta) = \sigma(\eta)(1 - \sigma(\eta))$, $[-\infty, +\infty] \rightarrow [0, 1]$ scale-invariant
- Log-odds/logit linear model: $\log \frac{p(y=1|x)}{1-p(y=1|x)} = \log \frac{p(y=1|x)}{p(y=0|x)} = x^T w + w_0$
- Predictions: $p(y = 1 | x) = \sigma(x^T w + w_0)$, $p(y = 0 | x) = 1 - \sigma(x^T w + w_0)$;
- MLE: Maximum Log-Likelihood, $w_* = \operatorname{argmax}_w \prod_{n=1}^N p(y_n | x_n | w) = \operatorname{argmax}_w \prod_{n=1}^N \sigma(x_n^T w) y_n (1 - \sigma(x_n^T w))^{1-y_n}$, equivalent to Logistic Loss
- Logistic loss: negative log-likelihood (NLL), convex surrogate of the 0-1 loss, $w_* = \operatorname{argmin}_w \sum_{n=1}^N -\log(p(y_n | w)) = \operatorname{argmin}_w \frac{1}{N} \sum_{n=1}^N -y_n x_n^T w + \log(1 + e^{x_n^T w})$
- Gradient: $\nabla L(w) = \frac{1}{N} \sum_{n=1}^N (\sigma(x_n^T w) - y_n) x_n = \frac{1}{N} X^T (\sigma(Xw) - y)$.
- Logistic loss penalizes deviation in the "wrong" direction more than the "correct" one
- Linearly separable data: MLE is unbounded so $\|w\| \rightarrow \infty$ increases likelihood (fix with ℓ_2 regularization). But it models a linear decision boundary
- No closed-form solution when minimizing negative log-likelihood
- SGD: updates gradient at all steps, the loss never goes to exact 0
- Newton's method: $w_{t+1} = w_t - \eta_t \nabla^2 L(w_t)^{-1} \nabla L(w_t)$ ($O(ND^2 + D^3)/\text{iter}$)

Support Vector Machines (SVM)

- Margin of hyperplane: $\min_{n \leq N} |w^T x_n|$. Distance from hyperplane: $|w^T x| / \|w\|_2$
- Max-Margin hyperplane: maximizes margin so if the training set slightly changes misclassifications will stay low $\max_w \|w\| = 1 (\min_{n \leq N} |w^T x_n|)$ s.t. $\forall n, y_n x_n^T w \geq 0$
- Hard-SVM: assume the dataset is linearly separable, $M = \frac{1}{\|w\|}$, $w' = \frac{1}{M} w$, $\min_w \frac{1}{2} \|w\|^2$ s.t. $\forall n, y_n x_n^T (w/M) \geq M \Rightarrow y_n x_n^T w' \geq 1$
- Soft-SVM: allow constraint violations $\min_{w, \xi} \frac{\lambda}{2} \|w\|^2 + \frac{1}{N} \sum_{n=1}^N \xi_n$ s.t. $\forall n, y_n x_n^T w \geq 1 - \xi_n$, simplified to $\min_w \frac{\lambda}{2} \|w\|^2 + \frac{1}{N} \sum_{n=1}^N [1 - y_n x_n^T w]_+$ with $[a]_+ = \max\{0, a\}$. Two terms balance regularization (variance) and accuracy (bias), w^* obtained with stochastic subgradient method
- Margin-based convex losses surrogate of 0-1 loss upper bounded by it ($\eta = yx^T w$):
 - Quadratic Loss (MSE): $\text{MSE}(\eta) = (1 - \eta)^2$, penalizes any deviation from 1
 - Logistic Loss: $\text{Logistic}(\eta) = \frac{\log(1+e^{-\eta})}{\log(2)}$, asymmetric cost; always penalizes.
 - Hinge Loss: $\text{Hinge}(\eta) = [1 - \eta]_+$, penalizes low-confidence predictions.
- Convex duality: $G(w, \alpha)$ s.t. $\min_w L(w) = \min_w \max_{\alpha} G(w, \alpha)$ (primal) and $\max_{\alpha} \min_w G(w, \alpha)$ (dual). Always $\max_{\alpha} \min_w G \leq \min_w \max_{\alpha} G$; equality if G convex in w , concave in α . Dual often easier; α typically sparse (support vectors).
- SVM w/ duality: $[z]_+ = \max_{\alpha \in [0, 1]} \alpha z$, $\min_w \max_{\alpha \in [0, 1]} N \left(\frac{1}{N} \sum_{n=1}^N \alpha_n (1 - y_n x_n^T w) + \frac{\lambda}{2} \|w\|^2 \right)$ convex in w and concave in α . Set $\nabla_w G = 0 \Rightarrow w(\alpha) = \frac{1}{\lambda N} \sum_{n=1}^N \alpha_n y_n x_n = \frac{1}{\lambda N} X^T \text{diag}(y) \alpha$. So: $\min_w L(w) = \max_{\alpha \in [0, 1]^N} \frac{1}{2\lambda N^2} \alpha^T \text{diag}(y) K \text{diag}(y) \alpha$, $K = XX^T$ (indep. of D ; solve via QP/coord ascent).

kNN and curse of dimensionality

- kNN: local averaging method, based on the values of labelled neighbours, no training, just query $O(ND)$. Small k : $\rightarrow$ Low bias & High variance (complex boundaries)
- Curse of Dimensionality: high $d \rightarrow$ Training set covers less space and points are far apart, $P(x \in \text{box}) = r^d = \alpha$. If $\alpha = 0.01$: $d = 10 \rightarrow r = 0.63$; $d = 100 \rightarrow r = 0.95$
- Generalization Bound for 1-NN: Bayes classifier f_* with $\eta(x) \geq 1/2$ with $\eta(x) = p(1 | x)$, Bayes risk $L(f_*) = \mathbb{E}_X [\min\{\eta(X), 1 - \eta(X)\}]$; assuming nearby points have the same label $\mathbb{E}_{S_{\text{train}}} [L(f_{S_{\text{train}}})] \leq 2L(f_*) + 4c\sqrt{dN} \frac{1}{\sqrt{2}} + 1$, last term is 0 if d constant and $n \rightarrow \infty$, increasing d and keep same err $\rightarrow$ increasing N exponentially
- Bound on Geometric Term: partition X into boxes $\{\text{Box}_k\}$, $p_k = \mathbb{P}(X \in \text{Box}_k)$. Then $\mathbb{P}(\exists x_i \in S_{\text{train}} : x_i \in \text{Box}_k | X \in \text{Box}_k) = 1 - (1 - p_k)^N$; if box side length ϵ , worst-case within-box distance $\|X - x_i\| \leq \sqrt{D} \epsilon$, distances between points $\rightarrow \infty$

Kernel Ridge Regression and the Kernel Trick

- Alt. Ridge Regr.: $w_* = X^T (XX^T + N\lambda I_N)^{-1} y$; $O(N^3 + dN^2)$ vs $O(d^3 + Nd^2)$. Regularization term lifts eigenvalues by $N\lambda$, so we always can invert the whole term
- Representer theorem: for any loss l , minimizer of $\frac{1}{N} \sum_{n=1}^N l(x_n^T w, y_n) + \frac{\lambda}{2} \|w\|^2$ satisfies $w_* = X^T \alpha_*$ (optimal solution in $\text{span}\{x_1, \dots, x_N\}$) with $\alpha_* \in \mathbb{R}^N$.
- Kernelized ridge (dual): $\alpha_* = \operatorname{argmin}_{\alpha} \frac{1}{2} \alpha^T \left(\frac{1}{N} XX^T + \lambda I_N \right) \alpha - \frac{1}{N} \alpha^T y$
- Feature extractor: map $\phi: \mathbb{R}^d \rightarrow \mathbb{R}^{\hat{d}}$ (often $\hat{d} \gg d$) to make data linearly separable; explicit dot-products in $\hat{d}$ are expensive.
- Kernel trick: pick $\kappa(x, x') = \phi(x)^T \phi(x')$ to achieve the $\mathbb{R}^{\hat{d}}$ dot-products without computing ϕ , also measures similarity of points, saves a lot on computation.
- Prediction: $y = \phi(x)^T w_* = \sum_{n=1}^N \kappa(x, x_n) \alpha_n$ (linear in $\mathbb{R}^{\hat{d}}$, nonlinear in $\mathbb{R}^d$).
- Quadratic kernel (scalar x): $\kappa(x, x') = (xx')^2$; Feature map: $\phi(x) = x^2$
- Polynomial kernel: $\kappa(x, x') = (x^T x')^2$ for $x, x' \in \mathbb{R}^3$; Feature map: $\phi(x) = [x_1^2, x_2^2, x_3^2, \sqrt{2}x_1x_2, \sqrt{2}x_1x_3, \sqrt{2}x_2x_3]^T \in \mathbb{R}^6$
- RBF kernel: $\kappa(x, x') = \exp(-\gamma \|x - x'\|_2^2)$ ($\gamma > 0$; implicitly infinite-dimensional ϕ).
- If κ_1, κ_2 kernels then $\alpha\kappa_1 + \beta\kappa_2$ ($\alpha, \beta \geq 0$) and $\kappa_1\kappa_2$ are kernels. Mercer: κ kernel iff symmetric and kernel matrix $\forall N \geq 0, \forall (x_n)_{n=1}^N, K = (\kappa(x_i, x_j))_{i,j=1}^N \succeq 0, \exists \phi$

Neural Networks

- Double descent: test risk decreases as # params increases, spikes near interpolation/memorization, then decreases again when highly overparametrized $\rightarrow$ NNs can generalize (extra degrees of freedom *wiggle room* beyond fitting train points).
- Neural Network (NN): $f_{NN}(x) = f(x)^T w$ (learned representation $f(x)$ + linear prediction w), non convex problem, $W^{(1)} \in \mathbb{R}^{d \times K}$, $W^{(l)} \in \mathbb{R}^{K \times K}$, $W^{(L+1)} \in \mathbb{R}^K$
- L -layer MLP: $x^{(l)} = \phi(W^{(l)T} x^{(l-1)} + b^{(l)})$; output $h(x) = W^{(L+1)T} x^{(L)} + b^{(L+1)}$; params $\approx O(K^2 L)$. SGD weight update: $(w_{i,j}^{(l)})_{t+1} = (w_{i,j}^{(l)})_t - \eta \frac{\partial \mathcal{L}_n}{\partial w_{i,j}^{(l)}}$
- Loss/output: regression $\frac{1}{2} (h(x) - y)^2$; binary $\log(1 + e^{-y h(x)})$ (logistic loss) or $\hat{y} = \operatorname{sign}(h)$; multiclass $-\log \sum_i \frac{e^{h_{ij}}}{\sum_i e^{h_{ij}}} = -h_y + \log \left(\sum_i e^{h_{ij}} \right)$ (Softmax Cross-Entropy w/ one-hot encoded y = numerically stable and convex log-sum-exp).
- Activation functions: if nonlinear it enables nonlinear models, sigmoid σ ; ReLU $[x]_+$ (applied element-wise); GELU $x \Phi(x) \approx x \sigma(1.702x)$ (differentiable ReLU at 0).
- Barron UAT: if $\int_{\mathbb{R}^d} \|w\| |f(w)| dw \leq C$, then $\exists f_N(x) = \sum_{j=1}^N c_j \phi(w_j^T x + b_j) + c_0$ s.t. $\int \|x\| \leq_r (f(x) - f_N(x))^2 dx \leq \frac{(2Cr)^2}{n}$, all smooth functions f can be approx. by a one-hidden-layer NN with error $O(1/n)$, high representational power. Sums of many simple (sigmoid/ReLU) units approximate complex functions (like Riemann integrals); L_2 averages error (can hide peak errors), L_{∞} controls worst-case errors.
- Piecewise-linear (PWL): any 1D PWL $q(x)$ can be written as affine + sum of shifted ReLUs, $q(x) = ax + b + \sum_{i=1}^n \alpha_i(x - \beta_i)_+$ (represented by 1-hidden-layer ReLU nets)
- Backprop: forward pass ($O(K^2 L)$) saves $x^{(l)}, z^{(l)}$. Error signal $\delta_j^{(l)} := \frac{\partial \mathcal{L}}{\partial z_j^{(l)}}$ with init $\delta^{(L+1)} = z^{(L+1)} - y$ and recursion $\delta^{(l)} = (W^{(l+1)} \delta^{(l+1)}) \odot \phi'(z^{(l)})$; gradients $\frac{\partial \mathcal{L}}{\partial w_{i,j}^{(l)}} = \delta_j^{(l)} x_i^{(l-1)}$, $\frac{\partial \mathcal{L}}{\partial b_j^{(l)}} = \delta_j^{(l)}$; compute all grads in $O(2K^2 L)$ (backward pass)
- Parameter Initialization: if init all weights and biases to 0 then all neurons in a layer act the same, chain rule multiplies many terms $\Rightarrow$ exploding/vanishing gradients; control layerwise variance (He init for ReLU: if width K , sample $w_i^{(l)} \sim \mathcal{N}(0, \sigma^2 I)$ with $\sigma^2 = 2/K$ so $\text{Var}[z^{(l)}] \approx 1$ across layers (input $\sigma^2 =$ output σ^2 for all layers)).
- BatchNorm: per feature (per neuron) normalize pre-activations in a mini-batch: $\hat{z}_n^{(l)} = \frac{z_n^{(l)} - \mu_B^{(l)}}{\sqrt{(\sigma_B^{(l)})^2 + \epsilon}}$, then $\hat{z}_n^{(l)} = \gamma^{(l)} \odot \hat{z}_n^{(l)} + \beta^{(l)}$ (learnable scale/shift so NN has the freedom to undo it); bias before BN is redundant (scale-invariant); regularization effect; inference uses fixed or running mean and variance.
- LayerNorm: normalize per sample across features: $\hat{z}_n^{(l)} = \frac{z_n^{(l)} - \mu_n^{(l)}}{\sqrt{(\sigma_n^{(l)})^2 + \epsilon}}$, then $\hat{z}_n^{(l)} = \gamma^{(l)} \odot \hat{z}_n^{(l)} + \beta^{(l)}$; same for train and inference; common in transformers; quicker and more stable training.

Convolutional Nets

- CNNs: less params than Fully-Connected NNs and capture complete spatial structure; convolution (+ nonlinearity) + pooling blocks then Fully-Connected classifier.
- Convolution: $x_n^{(l)} = \sum_{k,l=0}^{K-1} f_{k,l} x_{nS+k,mS+l}^{(l-1)}$. Weight sharing (same filter at every position) + local connectivity $\Rightarrow$ translation equivariance. Hyperparams: kernel/filter size K , stride S , padding P ; zero-padding helps borders and can preserve spatial size. Multi-channel conv: one filter per channel, sum across channels + bias.
- Pooling (max/avg): downsampling with hyperparams kernel/stride, no learnable ones
- Receptive field R_l grows with depth, early layers learn edges/colors, deeper ones learn objects. $W_{\text{out}} = \left\lfloor \frac{W_{\text{in}} + 2P - K}{S} \right\rfloor + 1$; $R_l = 1 + \sum_{j=1}^l ((K_j - 1) \times \prod_{i=1}^{j-1} S_i)$. Region size in layer l to make a pixel of final output: $r_l = S^{(l)} \cdot (r_{l+1} - 1) + K^{(l)}$, $r_{L+1} = 1$
- Backprop with weight sharing: sum gradients over all edges that share a weight.
- Residual nets: use skip connections $Y = R(X) + X$ (residual branch R) to mitigate vanishing gradients (gradient can flow through identity path by setting $R \approx 0$).
- Weight decay (ℓ_2 reg): $\min \mathcal{L} + \frac{\lambda}{2} \sum_i \|W^{(i)}\|_2^2$ (typically not on biases). GD update becomes $W = (1 - \eta\lambda)W - \eta \nabla \mathcal{L}$ (when the gradient vanishes weights go to 0 quickly); with BatchNorm, we get scale-invariance so we need other types of regularization.
- Dropout: during training keep each unit in layer l with prob. $p^{(l)}$ and train a random subnetwork; at test time use all units but scale activations by $p^{(l)}$ (or scale weights by $1/p^{(l)}$ during training). Regularizes but adds noise $\Rightarrow$ needs more iterations.

Ethics and Fairness in ML

- Avoid unfair discrimination and Stereotypes, understand Provenance of data.
- True error is an average criterion, ML models might achieve low error by minimizing for majorities and ignoring minorities.
- Removing a sensitive feature A may not remove its information if other features correlate with A .
- Fairness Criteria: score $R = f(X)$ where f is a model, outcome $D = 1_{R > t}$ with threshold t . A membership in protected group, Y real outcome; Independence: $R \perp A$ (equal acceptance rate); Separation: $R \perp A | Y$ (equal error rates, applied post-hoc); Sufficiency: $Y \perp A | R$ (calibration by group, A is irrelevant to predict Y using R).
- Calibration: $P(Y = 1 | R = r) = r$, by group if $P(Y = 1 | R = r, A = a) = r$
- Gaussian Distribution: $f(x; \mu, \sigma^2) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp \left(-\frac{(x - \mu)^2}{2\sigma^2} \right)$

Adversarial Machine Learning

- Data augmentation: make new training samples by transformation (regularization)
- Adversarial examples: small perturbations $\hat{x}$ causing misclassification confidently.
- Adversarial risk: $R_{\epsilon}(f) = \mathbb{E}[\max_{\|x-x_0\| \leq \epsilon} \ell(f(y(\hat{x})))]$ (g is the output of the NN before classification).
- White-box: taken wrt inputs (not the gradient, we change inputs to maximize loss) $\max_{\|x-x_0\| \leq \epsilon} \ell(y(g(\hat{x}))) = \ell(y(g(\hat{x}))) + \max_{\|\delta\| \leq \epsilon} \nabla x \ell(y(g(x)))^T \delta$ (Linearization) with $\delta = \hat{x} - x$. Closed-form 1-step attacks: ℓ_{∞} FGSM $\hat{x} = x + \epsilon \text{sign}(\nabla_x \ell(y(g(x))))$; ℓ_2 : $\hat{x} = x + \epsilon \nabla x \ell / \|\nabla_x \ell\|_2$. Iterative attacks via PGD: $\delta^{t+1} = \Pi_{B_p(\epsilon)}(\delta^t + \alpha \nabla_x \ell(y(g(x + \delta^t))))$ (for ℓ_{∞} often use sign). ℓ_2 +hinge loss=primal Soft-Margin SVM
- Black-box: score-based (we know g) or decision-based (we know predicted class $f(x)$) query attacks; finite-difference grad approximation $\nabla g(x) \approx \sum_i \frac{g(x + \alpha e_i) - g(x)}{\alpha} e_i$.
- Transfer attacks: use surrogate model to prepare White-box attacks in Black-box scenarios; can become model stealing if we can query the real model.
- Adversarial examples we produce should not be like the ones models commonly encounter during training (e.g. different camera angles, photographic distortions)
- Adversarial training: $\min_{\theta} \frac{1}{N} \sum_n \max_{\hat{x}} \|\hat{x} - x_n\| \leq \epsilon \ell(y_n g_{\theta}(\hat{x}))$ with inner maximizer $\hat{x}_n^*$ approximated by PGD; update θ with gradients on adversarial points $\frac{1}{N} \sum_n \nabla_{\theta} \ell(y_n g_{\theta}(\hat{x}_n^*))$. Robust and more interpretable gradients, but takes more PGD steps and has robustness/accuracy tradeoff (too large ϵ hurts clean accuracy). In comparison FGSM (ℓ_{∞}) only takes an extra forward+backward pass.

Unsupervised Learning

- Learn from unlabeled data $D = \{x_n\}_{n=1}^N$; representations (PCA/embeddings/matrix fact.), density/generative models (GMM/LLM), clustering/anomaly detection, EDA
- PCA is invariant to translation but not scaling, so apply stdscaling!
- Clustering: group similar objects, graph-based (partitioning), hierarchical (building a tree), model-based (maintain cluster models and infer membership K-means, GMMs).
- K-means: assign points to clusters with nearest mean, Within-Cluster Sum of Squares monotonically decreasing objective: $\min_{z, \mu} \sum_{n=1}^N \sum_{k=1}^K z_{n,k} \|x_n - \mu_k\|^2$, $z_{n,k} \in \{0, 1\}$, $\sum_k z_{n,k} = 1$. Non-convex & NP-hard, algorithm assigns nearest centroid to all points then update $\mu_k = \frac{\sum_n z_{n,k} x_n}{\sum_n z_{n,k}}$; coordinate descent $O(NKd)$, finite steps convergence, no lower bound. Choose K via domain knowledge / elbow method. Fails on non-spherical/overlapping clusters and outliers (drag centroids away from dense areas). Matrix Factorization $\sum_{n=1}^N \sum_{k=1}^K z_{n,k} \|x_n - \mu_k\|_2^2 = \|X^T - MZ^T\|_{Frob}^2$
- K-means++: better init than random (avoids poor local minima, faster convergence, lower objective). Pick μ_1 from dataset; define $D(x_n) = \min_{j \leq m} \|x_n - \mu_j\|_2^2$; sample next centroid with $\mathbb{P}(x_n) = \frac{D(x_n)}{\sum_n D(x_n)}$; repeat until K centroids, then run K-means.
- Probabilistic view: maximizing likelihood of isotropic Gaussians $\Leftrightarrow$ minimizing WCSS.

Expectation Maximization (EM) Algorithm

- GMM: pick clusters $Z \sim \text{Mult}(p)$, then sample $x \mid (Z = k) \sim \mathcal{N}(\mu_k, \Sigma_k)$ (x member of cluster k), $p(x) = \sum_{k=1}^K \pi_k \mathcal{N}(x \mid \mu_k, \Sigma_k)$, params $\theta = \{(\pi_k, \mu_k, \Sigma_k)\}_{k=1}^K$
- GMM MLE: $\mathcal{L}(\theta) = \sum_{n=1}^N \log \sum_{k=1}^K \pi_k \mathcal{N}(x_n \mid \mu_k, \Sigma_k) = \sum_{n=1}^N \log \sum_{k=1}^K p(x_n \mid \mu_k, \Sigma_k) p(x_n, \theta)$ (log-sum is hard; nonconvex; permutation of clusters gives same likelihood (model is not identifiable); can be unbounded). If labels were known $\rightarrow$ closed-form MLE (Hard-EM M-step): $\hat{\pi}_k = \frac{1}{N} \sum_n \mathbf{1}_{z_n=k}$, $\hat{\mu}_k = \frac{\sum_n \mathbf{1}_{z_n=k} x_n}{\sum_n \mathbf{1}_{z_n=k}}$, $\hat{\Sigma}_k = \frac{1}{\sum_n \mathbf{1}_{z_n=k}} \sum_n \mathbf{1}_{z_n=k} (x_n - \hat{\mu}_k)(x_n - \hat{\mu}_k)^T$.
- Hard-EM: assign hard labels $z_n = \arg \max_k \pi_k \mathcal{N}(x_n \mid \mu_k, \Sigma_k)$ then do MLE M-step
- K-means as special case of hard-EM: optimize only μ_k with fixed uniform weights $\pi_k = 1/K$ and fixed identical spherical covariances $\Sigma_k = \sigma^2 I$.
- Soft-EM: E-step: $q_{n,k} = p(z_n = k \mid x_n, \theta) = \frac{\pi_k \mathcal{N}(x_n \mid \mu_k, \Sigma_k)}{\sum_j \pi_j \mathcal{N}(x_n \mid \mu_j, \Sigma_j)}$ (probabilities of belonging to each cluster, captures uncertainty). M-step: $\pi_k = \frac{1}{N} \sum_n q_{n,k}$, $\mu_k = \frac{\sum_n q_{n,k} x_n}{\sum_n q_{n,k}}$, $\Sigma_k = \frac{1}{\sum_n q_{n,k}} \sum_n q_{n,k} (x_n - \mu_k)(x_n - \mu_k)^T$. EM steps use params that depend on each other recursively $\rightarrow$ cannot use closed-form solution. Use coordinate (monotonic) ascent: E-step sets $q^{(t+1)}(Z) = \arg \max_q \mathcal{L}(q, \theta^{(t)}) = \arg \max_q \mathbb{E}_{Z \sim q} \left[\log \left(\frac{p(X, Z | \theta^{(t)})}{q(Z)} \right) \right] = p(Z \mid X, \theta^{(t)})$; M-step $\theta^{(t+1)} = \arg \max_{\theta} \mathcal{L}(q^{(t+1)}, \theta) = \arg \max_{\theta} \mathbb{E}_{Z \sim q^{(t+1)}} [\log p(X, Z \mid \theta)] = \arg \max_{\theta} \mathbb{E}_{Z \sim q^{(t+1)}} [\log p(Z \mid \theta) p(X \mid Z, \theta)]$. $\mathcal{L}(\theta^{(t+1)}) \geq \mathcal{L}(\theta^{(t)})$, but can get stuck in local optima (not convex) and is not guaranteed to reach the global MLE.
- EM lower bound: introduce any distribution $q(Z)$ (the labels we never had), incomplete likelihood: $\mathcal{L}(\theta) = \mathcal{L}(q, \theta) + \text{KL}(q \parallel p(Z \mid X, \theta))$ upper bound on $\mathcal{L}(q, \theta) = \mathbb{E}_{Z \sim q} [\log p(X, Z \mid \theta)] + H(q)$ (KL distance between $\mathcal{L}(\theta)$, $\mathcal{L}(q, \theta)$)
- Initialization: π_k uniform, μ_k K-means++, Σ_k from empirical feature σ^2 , important!
- Pick K with Cross-Val against validation set log-likelihood. Clusters can overlap!
- Anomaly detection: after fitting a GMM, score new points by likelihood $p(x)$

Alternating Least-Squares

- No-missing simplification: two steps coordinate descent, $Z^T = (W^T W + \lambda_z I)^{-1} W^T X$, $W^T = (Z^T Z + \lambda_w I)^{-1} Z^T X^T$ (ridge).

Bias-Variance Decomposition

- Data Model: $y = f(x) + \epsilon$, $f(x)$ the true model, $\epsilon \sim \mathcal{D}_{\epsilon}$ (i.i.d. noise, $\mathbb{E}[\epsilon] = 0$)
- Bias: $(f(x_0) - \mathbb{E}_S[f_S(x_0)])^2$, how far the expected prediction is from the true value $f(x_0)$. Low model complexity $\rightarrow$ high bias, high complexity $\rightarrow$ low bias.
- Variance: $\mathbb{E}_S[(f_S(x_0) - \mathbb{E}_S[f_S(x_0)])^2]$, measures the variability of predictions at x_0 across different training sets. High model complexity $\rightarrow$ high variance.
- Error Decomposition: at x_0 is $L(f_S) = \mathbb{E}_{\epsilon} \sim \mathcal{D}_{\epsilon} [(f(x_0) + \epsilon - f_S(x_0))^2]$
- Bias-Variance Decomposition: $\mathbb{E}_{S, \epsilon} [(f(x_0) + \epsilon - f_S(x_0))^2] = \text{Var}[\epsilon] + (f(x_0) - \mathbb{E}_S[f_S(x_0)])^2 + \mathbb{E}_S[(f_S(x_0) - \mathbb{E}_S[f_S(x_0)])^2] = \text{Noise} + \text{Bias}^2 + \text{Variance}$

Transformers

- Transformer: learned function that transforms input into contextualized representations by mixing information across tokens via self-attention. Text and images are flattened and mapped to real-valued vectors. Transformer block composed of (LayerNorm + SelfAtt) with a SkipConn and then (LayerNorm + MLP) with a SkipConn. $O(L(T^2 K + T K^2))$ with L layers of width K and sequence length T
- Tokenization: split text in tokens with integer ID based on a vocabulary V
- BPE makes vocab, starting with a token per char, merging until desired vocab. size
- Token embeddings matrix $\in \mathbb{R}^{|V| \times D}$: map token ID to a $\mathbb{R}^d$ vector (input $X \in \mathbb{R}^{T \times D}$), mappings start random, learned via backprop during Transformer's training
- softmax $_{\mathcal{Z}}(z)_i = \frac{e^{z_i}}{\sum_j e^{z_j}}$, $\mathbb{R}^d \rightarrow [0, 1]^d$, d dimension of $z \in \mathbb{R}^d$, scale-shift invariant
- Label smoothing: $\hat{y} = (1 - \epsilon)y + \epsilon / D$ (replaces hard one-hot targets with softer ones; reduces overconfidence & huge logits in softmax+CE, improves numerical stability).
- Self-attention: transforms contextualized representations using learned input-dependent weighted averages, $Q = XW_Q$, $K = XW_K$, $V = XW_V$; $P = \text{softmax}(QK^T / \sqrt{D_K})$ (row-wise, weighting coeffs $\in [0, 1]$); output $Z = PV$. $p_{i,j}$ is based on how similar q_i, k_j are (dot product), normalize by $\sqrt{D_K}$ (stabilizes gradients for faster convergence). Positional embeddings W_{pos} learned via backprop are added to input as attention is order-agnostic (map position $\{1, \dots, T\} \rightarrow$ feature vector $\mathbb{R}^D$). Captures long-range dependencies within the context window. MLP layer applies same transformation to all tokens independently, $\text{MLP}(X) = \phi(XW_1 + b_1)W_2 + b_2$ (W_1, W_2, b_1, b_2 learned via backprop). $O(T^2)$, parallelizable
- Multi-head: many parallel attention patterns (like multiple channels in CNNs), compute Z_h per head, concatenate $Z = [Z_1, \dots, Z_H]W_O$, W_O learned via backprop
- Causal masks for decoders: sets $p_{i,j} = 0$ for $j > i$ (or add $-\infty$ above diagonal before softmax) to not predict based on following tokens.
- Used in Encoders (tagging (per-token classification), seq. classification), Decoders (LLMs) and Encoder-decoders (seq2seq, translation). Everything can be tokenized!

LLMs & Text Representation Learning

- Co-occurrence counts: $N = (n_{ij}) = \#\{\text{contexts with } w_i \text{ and } w_j\}$ sparse matrix, need vocabulary & context definition, often uses log counts, perfect for matrix fact.
- GloVe: weighted MF with $f_{dn} = \min(1, (n_{dn}/n_{max})^\alpha)$, train via SGD or ALS.
- Skip-gram/word2Vec: LogReg separates real (w_d, w_n) (window size ≈ 5) and fake (negative sampling, distinguish noise from data) maximize dot products for real pairs
- Words w_d with opposite meanings can have very similar context words w_n .
- Sentence/doc repr.: supervised (CNN/transformers/MF); fastText uses Bag of Words $x_n \in \mathbb{R}^{|V|}$ with 2 layers (learned embeddings W + linear classifier Z), $\min_{W, Z} \mathcal{L}(W, Z) = \sum_{s,n} f(y_n(WZ^T x_n))$ (not jointly convex). Unsupervised: averaging word vectors (baseline); train word vectors so averaging works (better); current: training with sentences as inputs to capture context (no single words, no BoW).
- Language models: self-supervised next-token prediction (massive general multi-task); multi-class (softmax over vocab.) or binary (try many word pairs with new words, pick the best one deemed real); factorization $p(x_1, T) = \prod_t p(x_t \mid x_{<t})$; Autoregressive inference (forward pass, pick next token by sample or argmax, append, repeat).
- Pretraining: general-purpose model from massive text; data sources (CommonCrawl, code, curated, synthetic); challenges: quality/diversity + trillion-token efficiency.
- GPT-1/2/3 (decoder-only), larger models improve (finetuning becomes optional, in-context learning), number of params and training set size increases.
- In-context learning: zero/one/few-shot task completion via examples in the prompt.
- Chain of Thought: showing examples of reasoning steps helps with complex tasks
- Posttraining: pretrained models may not follow user intent/policies; Supervised Finetuning is expensive and penalizes all mistakes equally and noisy (human labelling).
- Collect pairwise comparisons (x_{prompt}, x_a, x_b) , train reward model $R_{\phi}: X_{prompt} \times X_{response} \rightarrow \mathbb{R}$ with loss $\mathbb{E}[-\log \sigma(R_{\phi}(x_{prompt}, x_a) - R_{\phi}(x_{prompt}, x_b))]$.
- RLHF: finetune the policy π_{θ} (pretrained LLM) to maximize $J(\theta) = \mathbb{E}_{x_{prompt} \sim P_{data}, x \sim \phi_{\theta}} [R_{\phi}(x_{prompt}, x)]$; policy gradient $\nabla_{\theta} J = \mathbb{E}[R_{\phi}(x_{prompt}, x) \nabla_{\theta} \log \pi_{\theta}(x \mid x_{prompt})]$ (log-derivative trick, because we cannot not prompt the derivative of the expected reward defined over the policy distribution ϕ_{θ} since it depends on the params θ we are optimizing).
- PPO: constrain/penalize updates to $\nabla_{\theta} J$ so that pretraining is not forgotten.
- Evaluation: loss/acc not comparable across tokenizers; use open-ended challenges + benchmarks (MMLU/ARC) + model/human anenas.

Matrix Factorization

- Low-rank model: $X \approx WZ^T$ with item factors $W \in \mathbb{R}^{D \times K}$, user factors $Z \in \mathbb{R}^{N \times K}$.
- Objective: $\min_{W, Z} \frac{1}{2} \sum_{(d,n) \in \Omega} \|x_{dn} - (WZ^T)_{dn}\|^2 + \frac{\lambda_w}{2} \|W\|_F^2 + \frac{\lambda_z}{2} \|Z\|_F^2$.
- Not jointly convex / not identifiable; choose K too large $\rightarrow$ overfit.
- SGD on one observed (d, n) : with error $e_{dn} = x_{dn} - w_d^T z_n$, update uses $\nabla w_d = -e_{dn} z_n$, $\nabla z_n = -e_{dn} w_d$ (+ reg.).

Self-Supervised Learning

- Transfer learning: finetune pretrained model on a related task to the one it was trained for using few labels (they are expensive), good performance on the related task.
- Self-supervised: create labels from input x via $g: x \mapsto (x_{in}, x_{out})$, learn pretext task $f: x_{in} \rightarrow x_{out}$ (e.g. text = MLM / next-token; vision = rotation, colorization).
- Masked Language Modelling (MLM): mask tokens ([MASK]) and predict them, learns contextual representations for downstream tasks.
- BERT (encoder-only bidirectional transformers): input = 1–2 sentences with [CLS] and [SEP]; encoder attends to full sequence. Masking: replace $\approx 15\%$ tokens with [MASK] (sometimes random token or unchanged to reduce train-test shift for downstream tasks that don't do masking). Embeddings: token + positional + segment/sentence membership embeddings. Predict original tokens for masked positions in parallel independently (softmax cross-entropy); NSP via [CLS] predicts if sentence B follows A. Can be finetuned for sequence classification and sentiment analysis via [CLS], token classification (NER) via per-token outputs, retrieval tasks.
- Joint embedding Methods: encoder invariant to some transformations of data, two augmented views $(x_1, x_2) = g(x)$, learn encoder $f: X \rightarrow \mathbb{R}^d$ s.t. $f(x_1) \approx f(x_2)$; avoid collapse (f always predicting the same constant) via contrastive or regularized non-contrastive learning (e.g. BYOL w/ exponential moving average target encoder).
- Contrastive Learning: push away representations of unrelated views from different data points, x and a view of it x^+ , any negative view x^- we want $s(f(x), f(x^+)) > s(f(x), f(x^-))$ where s is the similarity of two embeddings
- SimCLR: sample two augmented views (x, x^+) + many negatives $\{x_i^-\}$; optimize $\text{InfoNCE } L(x) = -\log \frac{e^{s(f(x), f(x^+))}}{e^{s(f(x), f(x^+))} + \sum_{i=1}^N e^{s(f(x), f(x_i^-))}}$, $f(x) = f_2(f_1(x))$; encoder f_1 + projector f_2 (train both end-to-end; use f_1 downstream); cosine similarity with temperature τ : $s(e_1, e_2) = \frac{\langle e_1, e_2 \rangle}{\|e_1\|_2 \|e_2\|_2 \tau}$; strong data augmentations.
- CLIP: contrastive image-text; maximize cosine similarity of matched (image, caption) pairs, minimize mismatched pairs; capable of few-shot learning+zero-shot prompting.

Generative ML

- Discriminative models $p(y \mid x)$; Generative models model $p(x)$ (or $p(x \mid y)$)
- Explicit density models define a density/likelihood; Implicit density models generate realistic samples but can suffer from mode collapse (low diversity, e.g. GANs)
- Tractable density models: explicit, exact likelihood, Autoregressive (LLMs, slow sampling) or Normalizing Flows (efficient sampling/inference via invertible transformations, computationally intensive)
- Approximate models: explicit, VAE (optimize lower bound on likelihood, stable but blurrier samples), Energy-based or Diffusion (high quality audio/images but slow).
- GAN: generative adversarial network, generator $G(z) \in p_g$ with $z \sim p_z$ (noise) and discriminator $D(x) \in (0, 1)$ with $x \sim p_g$; differentiable two-player min-max game $\min_G \max_D \mathbb{E}_{x \sim p_g} [\log D(x)] + \mathbb{E}_{z \sim p_z} [\log(1 - D(G(z)))]$; first term is how good is D at recognizing real samples (can't be used for classification); second term is G wanting to fool D as much as possible while D does the opposite; optimum when $p_g = p_d$. Trained with Adam by Alternating-GAN (sample x and z , maximize D , sample z again, minimize G , repeat). Conditional if it conditions G, D on extra info (e.g. image, text or one-hot encoded class labels).
- Deep Convolutional GAN: for images, deconvolution layers that do image upsampling
- Diffusion models: progressively add noise via forward Markov chain $p(x_0:T) = p_0(x_0) \prod_{t=0}^{T-1} p(x_{t+1} \mid x_t)$, then generate samples by reversing it: $p(x_0:T) = p_T(x_T) \prod_{t=0}^{T-1} p(x_t \mid x_{t+1})$ (backward transition, iterative denoising).
- Ancestral sampling (score-based): start $X_T \sim p_{\text{ref}} \approx \mathcal{N}(0, 1)$ and for $t = T-1, \dots, 0$ sample $X_t \sim p(\cdot \mid X_{t+1})$. With Gaussian forward step $p(x_{t+1} \mid x_t) = \mathcal{N}(\alpha x_t, (1 - \alpha^2)I)$, reverse is approximated w/ Taylor expansion of $\log p(x_t)$ around x_{t+1} (because we have no way of knowing $p(x_t \mid x_{t+1})$ when going backward), $p(x_t \mid x_{t+1}) = p(x_{t+1} \mid x_t) \exp[\log p(x_t) - \log p(x_{t+1})] \approx \mathcal{N}(x_t; (2 - \alpha)x_{t+1} + (1 - \alpha^2)\nabla \log p(x_{t+1}), (1 - \alpha^2)I_d)$ where $\nabla \log p(x_{t+1})$ is the score (intractable as it requires calculating integral over every possible clean image x_0). Approximate it using a NN $s_{\theta}(t, x)$ (Denoising Score Matching) conditioned on X_0 with loss $\mathcal{L}(\theta) = \frac{1}{2} \mathbb{E}_{X_t} [\|\nabla \log p(X_t) - s_{\theta}(X_t)\|^2] = C_2 \frac{1}{2} \mathbb{E}_{X_0} \mathbb{E}_{X_t \mid X_0} [\|s_{\theta}(X_t) - \nabla \log p(X_t \mid X_0)\|^2]$; predict directly noise added $\hat{\xi}_{\theta}(t, X_t)$ with $s_{\theta}(t, X_t) = -\hat{\xi}_{\theta}(t, X_t) / \sqrt{1 - \alpha^{2t}} \rightarrow \mathcal{L}(\theta) = \frac{1}{2} \sum_{t=1}^T \mathbb{E}_{X_0} \mathbb{E}_{X_t \mid X_0} [\|\hat{\xi}_{\theta}(t, X_t) - \xi_t\|^2]$. Train DSM during forward transitions updating θ with SGD computed at a random t and $\xi_t \sim \mathcal{N}(0, I_d)$. Generate samples starting with $X_T \sim p_{\text{ref}} \approx \mathcal{N}(0, I)$ and iteratively denoising it: $X_t = (2 - \alpha)X_{t+1} + (1 - \alpha^2)s_{\theta}^*(t+1, X_{t+1}) + \sqrt{1 - \alpha^2}\xi_{t+1}$
- U-net architecture: DSM implemented with encoder-decoder (self-attention) + skip connections (ResNet). Time t embedded with sinusoidal positional encodings and injected into residual blocks (spatial addition or adaptive group normalization layers). Exponential Moving Average of parameters for stability: $\theta_{n+1} = (1 - \beta)\theta_n + \beta\theta_n$.
- Conditional diffusion: condition DSM on additional input (label/prompt Y) to learn $s_{\theta}(t, X_t; Y)$, embedded like time embeddings in U-net; objective becomes $\theta^* = \arg \min_{\theta} \frac{1}{2} \sum_{t=1}^T \mathbb{E}_{X_0, Y} \mathbb{E}_{X_t \mid X_0} [\|s_{\theta}(t, X_t^Y; Y) - \nabla \log p(X_t^Y \mid X_0^Y)\|^2]$

Lab Exercises

- $f(z) = \sum_{i=0}^d a_i z^i$ pos. coeffs, k_1 kernel, $k(x, x') = \sum_{i=0}^d a_i (k_1(x, x'))^i$ is a kernel
- Numerically Stable Softmax: $\exp(x_i - x_{max}) / \sum_j \exp(x_j - x_{max})$
- $H = -\sum_{i=1}^d y_i \ln(z_i)$; $\frac{\partial H}{\partial x_j} = \sum_{k=1}^d \frac{\partial H}{\partial z_k} \frac{\partial z_k}{\partial x_j}$; $\frac{\partial H}{\partial x_j} = \frac{\partial H}{\partial z_j} x_j + \sum_{k \neq j} \frac{\partial H}{\partial z_k} \frac{\partial z_k}{\partial x_j}$
- Minimize dot product between w, x ; x must point in opposite direction of $w, w/\|w\|_2$
- $\nabla_x (w^T A w) = (A + A^T)w$, if A was symmetric ($A = A^T$), it would just be $2Aw$